

KICKSTARTER

For creatorsLog in

Q Search projects, creators, and categories

s Crafts Dance Design Fashion Film Food Games Journalism Music Photography Publishing Technology Theater



InsightAI — The Cognitive AI, Owned. Not Subscribed.

InsightAI accumulates. Sovereign cognitive AI, persistent memory, Expert engines on your hardware.



Shahin Talaei

First created • 0 backed

More

\$923

pledged of \$17,653 goal

14

backers

 Software

 Toronto, Canada

Funding Unsuccessful

The project's funding goal was not reached on Fri, June 12 2026 4:15 AM UTC +06:00

[Campaign](#)[Rewards](#)[Creator](#)[FAQ](#)[Updates ⁰](#)[Comments ¹](#)[Community](#)

Story

We've built the most expensive amnesiacs in history.

Every AI you use today is brilliant for an hour and gone by morning. You pay twenty dollars a month to start over, every single day, with software that learns from you and keeps the lessons for itself.

This is the architecture we've all just accepted.

I'm not interested in accepting it.

Imagine opening an AI on a Tuesday morning. It already knows you're three weeks into the rewrite. That your daughter's recital is Friday. That you stopped using the word "leverage" six months ago because it sounded fake. It knows the project. It knows your voice. It knows what you've already tried and what you've already rejected.

You don't onboard it.

You resume.

That's what an AI is supposed to feel like. That's what you should have been getting all along.

InsightAI is a cognitive AI that runs on your hardware. Not in the cloud. Not in someone's data center. On the machine sitting in front of you.

It remembers every conversation you have with it across sessions, across days, across years. It routes your work through specialized

Funding period

May 13 2026 - Jun 12 2026
(30 days)

expert engines, each tuned for a different kind of thinking. It adapts to how you write, what you prefer, what you reject.

Pay once. Own it forever. Your data never leaves the room.

The real reason I'm building this

When the AI engines first came out, I knew immediately that something had shifted. Not because the model was good. It wasn't, not really, not yet but because for the first time I could talk to a machine that seemed to understand what I was saying.

So I started building. I used every system I could get my hands on — ChatGPT, Claude, Gemini, Grok — pushing each of them as far as they'd go. And every single one had the same wall: the moment I closed the tab, the system forgot I existed. Whatever we'd built together, whatever insight we'd reached, whatever it had learned about how I think, gone.

Two years in, I stopped thinking of memory as a feature.

I started thinking of it as a precondition.

You cannot build a mind that resets every Tuesday. You cannot evolve something that begins again at zero with every conversation. Whatever artificial general intelligence is going to be, when it arrives, it will have woken up inside a system that remembered yesterday.

InsightAI started as that bet. It's still that bet. It's a sovereign, self-hosted cognitive substrate built around the premise that memory and continuity aren't optional they're the floor. The minimum. The thing without which nothing else matters.

Will InsightAI become AGI? I don't know. I'm not promising you that. What I am promising is that it has the architectural shape that cloud AI doesn't and that the people who back this campaign will own their copy of it before that question gets answered.

I'd rather be wrong trying than right by accident.

What you're funding

The platform exists. It's running. It's documented. The landing page is live at insightai.life. The codebase is real — Next.js frontend, FastAPI backend, PostgreSQL and Redis under the hood, Llama 3.3 70B running locally.

What's missing is the final mile: a one-click installer for Mac, Windows, and Linux. Production hardening. A complete onboarding experience. Founding-member features that won't ship in the public release.

That's what your money builds.

Every Founder Edition Lifetime License is numbered. There will only ever be one N° 001.

The next decade of AI will be decided by a single question:

Who owns the intelligence?

If the answer is few companies, the future of thought routes through few buildings, behind few logins, on few sets of terms-of-service that can change tomorrow.

I don't want to live in that world.

I doubt you do either.

There is another way.

Pay once. Run it on your hardware. Watch it remember you. Watch it become yours.

Back this campaign and you don't just buy a product.

You take a position.

Risks and challenges

I'll be honest about what could go wrong, because the worst thing a Kickstarter founder can do is pretend the path is clear.

Timeline risk

I'm a solo technical founder. The platform is at roughly 90% production readiness today, but the final 10% — installer polish across three operating systems, onboarding, and production hardening — historically takes longer than founders predict. My realistic timeline is 4–6 months post-funding to ship. If I hit a wall I can't see from here, I will tell backers directly through campaign updates rather than quietly slip the date.

Technical risk

InsightAI runs locally, which means it depends on each backer's hardware. The platform is designed for modern machines (Apple Silicon, recent Nvidia/AMD GPUs, or sufficient CPU + RAM for CPU-only mode), but performance will vary. I will publish minimum and recommended hardware specs before fulfillment so no backer is surprised. Backers whose hardware doesn't meet specs will be offered a refund or a transferable license.

Model dependency risk

InsightAI runs on open-weight models — currently Llama 3.3 70B with a Llama 3.1 8B fallback. If Meta or another provider materially changes their license terms in a way that affects redistribution, I will migrate to whichever open-weight model best preserves the architecture. Open-weight alternatives are growing, not shrinking, and the platform is model-agnostic by design.

Solo founder risk

I am one person. If I'm hit by a bus, the project is at risk. To address this: the codebase is documented, the architecture is mapped, and post-funding I will deposit a release-of-source clause with a third party so that backers who paid for a lifetime license retain access to the source in the event the project cannot continue. This is a commitment, not a marketing line — I will publish the legal mechanism before fulfillment.

The bigger question

Some people will read the project page and worry that the AGI ambition means I'm overpromising what InsightAI will become. I want to address this directly: I am not promising that InsightAI will become AGI. I am promising that it has the architectural shape — persistent memory, expert specialization, behavioral learning, sovereign deployment — that I believe is required for AGI to emerge. Whether that emergence happens, and on what timeline, is genuinely unknown. What backers receive on day one is a working sovereign cognitive AI. Anything beyond that is a bet we're making together.

If something material changes, I will tell you in the update feed before you read it anywhere else. That's the contract.

[Learn about accountability on Kickstarter](#)

Use of AI

My project seeks funding for AI technology.

I plan to use AI-generated content in my project.

What parts of your project will use AI generated content? Please be as specific as possible.

1.VISUAL ASSETS (algorithmically generated) The Founder Edition Lifetime License card image was produced programmatically using the Python Pillow imaging library — a vector/raster code-rendered graphic, not output from a generative image model trained on third-party artwork. The design direction was developed in conversation with Claude; the rendering is mathematical (gradients, geometry, typography). No third-party copyrighted artwork was referenced, scraped, or reproduced. 2.ARCHITECTURE DIAGRAMS (programmatically generated SVG) The system architecture diagrams that may be shown on the project page are SVG/HTML rendered from my own codebase descriptions — original technical illustrations of my own architecture, not generated from external creative works.

Do you have the consent of owners of the works that were (or will be) used to produce the AI generated portion of your projects? Please explain.

FOR THE VISUAL ASSETS (algorithmically rendered) The Founder card and architecture diagrams were rendered programmatically using standard imaging and SVG libraries (Python Pillow, raw SVG). These are mathematical constructions of geometry, gradients, and typography — not generative outputs trained on artist data. No artist or photographer's work was referenced, sampled, or imitated in the production of these images. The fonts used (Poppins, by Indian Type Foundry) are open-source under the SIL Open Font License.

[Learn about AI policy on Kickstarter](#)

Environmental commitments

[Visit our Environmental Resources Center](#) to learn how Kickstarter encourages sustainable practices.

Long-lasting design

InsightAI is built for permanence, not obsolescence.

Three deliberate architectural choices extend the life of what backers receive far beyond a typical software product:

1. Self-hosted, no server dependency. The platform runs entirely on the user's own hardware. There is no central server that can be shut down, no cloud service that can be sunset, and no subscription that can lapse. The software a backer receives on day one will continue to function indefinitely, regardless of the company's future.
2. Open-weight model foundation. InsightAI is built on open-weight large language models (Llama 3.3 70B / 3.1 8B), not closed APIs. This means backers are never at the mercy of a vendor deprecating an API or revoking access. The models live on disk, locally, and run without external dependency.
3. Documented, transferable architecture. The codebase is thoroughly documented and the architecture is mapped in public diagrams. As part of my commitment to backers, I will deposit a release-of-source clause with a third party so that, in the unlikely event the project cannot continue, lifetime-license holders retain perpetual access to the source code and the right to maintain their own copies.

Compared to subscription cloud AI — which by design has a finite life tied to vendor business decisions — InsightAI

is engineered to outlive its creators. That is the literal opposite of planned obsolescence.

Sustainable Distribution

InsightAI is delivered as a 100% digital product. There is nothing to ship, nothing to package, and nothing to fly across an ocean.

Specifically:

- Zero physical packaging. No boxes, no plastic, no printed manuals, no shipping inserts.
- Zero transportation footprint. Backers download the software directly from a CDN; there is no warehouse, no fulfillment center, no last-mile delivery truck.
- Zero returns logistics. Software does not need to be returned, repaired, or replaced through reverse-logistics channels.
- Local-first inference. Once installed, InsightAI runs on hardware the user already owns. Unlike cloud AI — which routes every query through hyperscale data centers consuming significant power and water — local inference adds no incremental data-center load per use.

The most sustainable shipping route is the one that doesn't exist. By designing InsightAI as a sovereign, locally-run, digitally-delivered product, the campaign's environmental footprint per backer is bounded almost entirely by the energy used to download a single installer file.

For digital infrastructure I do control (the download CDN, the project website, supporting services), I prioritize providers with verifiable renewable-energy commitments.

Something else

Local-first AI architecture. Cloud AI inference at scale runs on hyperscale data centers consuming significant power

and water for cooling. InsightAI runs inference on hardware the user already owns and powers, eliminating the incremental data-center load per query. A user running InsightAI for a year typically uses less total energy than the data-center share of equivalent cloud-AI usage.

Additionally: 100% digital delivery (no packaging, no shipping, no physical fulfillment), and the platform is designed to run efficiently on existing consumer hardware rather than requiring users to purchase new equipment.

Questions about this project? [Check out the FAQ](#)

[Report this project to Kickstarter](#)

Similar projects to check out

[See more](#)





 Save our oceans. Take Wildtype mainstream.

Justin Kolbeck

 10 days left • 107% funded





 **Zen Tea Brew Infusions: Discover Delicious Rare Shizuoka Tea** 
Zen Tea Brew



 **Faire in the Field: Year Two!** 
Patrick Munroe-Smith



 **Laughing Gull Chocolates Take Flight** 
Lindsay
 10 days left • 81% funded



Discover

Arts

Comics & Illustration

About

About us

Our charter

Support

Help Center

Our Rules

More

Newsletters

Kickstarter Project Updates

KICKSTARTER

English

▼

\$ US Dollar (USD)

▼

