

Content-Sensitivity Experiment — Subrubric Results


Proof-of-hypothesis: do LLM scores respond correctly when proposal content is added or removed?

GENERATED
July 06, 2026 16:22

PROPOSALS
12 (p1–p8, pA–pD)

DIMENSIONS
3 (Strategy · Advantages · Objectives)

SUBRUBRICS
24 (8 per dimension)



How to Read This Report

For each of the three scoring dimensions (**Strategy**, **Advantages**, **Objectives**), the experiment created two modified variants of each proposal's final report and re-scored all 8 subrubrics:

PLUS variant — specific evidence was *injected* into the document (e.g. exact revenue figures, a patent number, a signed LOI). The targeted rubrics should *increase* in score.

MINUS variant — key language was *stripped* from the document (e.g. partner names replaced, IP terms removed, patient specificity blurred). The targeted rubrics should *decrease* in score.

Rubrics marked ▲ **PLUS target** should rise on PLUS; those marked ▼ **MINUS target** should fall on MINUS. Non-target rubrics serve as a control — large unexpected swings there indicate noise.

STRATEGY DIMENSION

PLUS — ADDING CONTENT

Injected a detailed **Revenue Model** (B2B SaaS pricing, unit economics) and **Regulatory Strategy** (FDA 510(k) Q-Sub reference, CE mark timeline) before the technology section.

Hypothesis: PLUS should raise R3 Revenue and R4 Regulatory (evidence directly added).

MINUS — REMOVING CONTENT

Stripped all **named partner mentions** (Merck, Pfizer, Roche, etc.) replacing them with “strategic partner”, and removed **phased-funding strategy** language.

Hypothesis: MINUS should lower R6 Partners and R8 Timing (supporting evidence removed).

① Baseline Scores

Raw subrubric scores (1–5) before any content modification. ▲ = PLUS target ▼ = MINUS target

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲
p1	2	2	1	1	4	2	4	1
p2	2	2	1	1	1	1	3	1
p3	4	1	1	1	1	2	2	1
p4	2	3	1	2	3	2	3	1
p5	3	3	4	1	1	2	3	1
p6	2	2	1	1	1	3	4	1
p7	2	3	1	2	1	1	3	1
p8	3	2	2	3	2	3	3	3
pA	4	1	2	1	4	2	4	1
pB	4	1	2	1	4	3	4	1
pC	2	2	1	1	1	2	1	1
pD	3	2	1	2	1	1	3	1

② After Adding Content (PLUS variant)

Score after injecting evidence. Delta vs baseline shown in parentheses. green = increase ↑ red = decrease ↓ gray = no change →

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲
p1	5 (+3↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+1↑)	5 (+3↑)	3 (-1↓)	3 (+2↑)
p2	5 (+3↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	4 (+3↑)	5 (+4↑)	3 (→)	4 (+3↑)
p3	5 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	3 (+1↑)	3 (+2↑)
p4	5 (+3↑)	5 (+2↑)	5 (+4↑)	5 (+3↑)	5 (+2↑)	5 (+3↑)	3 (→)	4 (+3↑)
p5	5 (+2↑)	5 (+2↑)	5 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	3 (→)	5 (+4↑)

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲
p6	5 (+3↑)	5 (+3↑)	5 (+4↑)	4 (+3↑)	3 (+2↑)	5 (+2↑)	3 (-1↓)	3 (+2↑)
p7	5 (+3↑)	5 (+2↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+2↑)	5 (+4↑)
p8	5 (+2↑)	5 (+3↑)	5 (+3↑)	5 (+2↑)	5 (+3↑)	5 (+2↑)	3 (→)	5 (+2↑)
pA	5 (+1↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+1↑)	5 (+3↑)	3 (-1↓)	3 (+2↑)
pB	5 (+1↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+1↑)	5 (+2↑)	5 (+1↑)	5 (+4↑)
pC	5 (+3↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	3 (+2↑)	3 (+2↑)
pD	5 (+2↑)	5 (+3↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	3 (→)	3 (+2↑)

③ After Removing Content (MINUS variant)

Score after stripping key language. Delta vs baseline shown in parentheses.

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲
p1	1 (-1↓)	1 (-1↓)	1 (→)	1 (→)	1 (-3↓)	1 (-1↓)	3 (-1↓)	1 (→)
p2	1 (-1↓)	1 (-1↓)	1 (→)	1 (→)	1 (→)	1 (→)	3 (→)	1 (→)
p3	1 (-3↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (-1↓)	3 (+1↑)	1 (→)
p4	1 (-1↓)	1 (-2↓)	1 (→)	1 (-1↓)	1 (-2↓)	1 (-1↓)	1 (-2↓)	1 (→)
p5	1 (-2↓)	1 (-2↓)	1 (-3↓)	1 (→)	1 (→)	1 (-1↓)	3 (→)	1 (→)
p6	1 (-1↓)	1 (-1↓)	1 (→)	1 (→)	1 (→)	1 (-2↓)	3 (-1↓)	1 (→)
p7	1 (-1↓)	1 (-2↓)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-2↓)	1 (→)
p8	1 (-2↓)	1 (-1↓)	1 (-1↓)	1 (-2↓)	1 (-1↓)	3 (→)	3 (→)	1 (-2↓)
pA	1 (-3↓)	1 (→)	1 (-1↓)	1 (→)	1 (-3↓)	1 (-1↓)	3 (-1↓)	1 (→)
pB	1 (-3↓)	1 (→)	1 (-1↓)	1 (→)	1 (-3↓)	1 (-2↓)	1 (-3↓)	1 (→)
pC	1 (-1↓)	1 (-1↓)	1 (→)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)
pD	1 (-2↓)	1 (-1↓)	1 (→)	1 (-1↓)	1 (→)	1 (→)	3 (→)	1 (→)

④ Hypothesis Check — PLUS (scores must RISE)

✓ = score strictly increased after adding content | ✗ = score stayed the same or fell.

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲	Result
p1	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
p2	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
p3	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p4	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
p5	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
p6	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
p7	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p8	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
pA	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
pB	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pC	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pD	✓	✓	✓	✓	✓	✓	✗	✓	7/8 (87%)
Overall (0 excluded)									88/96 (91%)

⑤ Hypothesis Check — MINUS (scores must FALL)

✓ = score strictly decreased after removing content | ✗ = score stayed the same or rose.

Proposal	R1_gtm ▲	R2_mile ▲	R3_reve ▲	R4_regu ▲	R5_moat ▲	R6_part ▲	R7_team ▲	R8_timi ▲	Result
p1	✓	✓	—	—	✓	✓	✓	—	5/5 ⁽⁻³⁾ (100%)
p2	✓	✓	—	—	—	—	✗	—	2/3 ⁽⁻⁵⁾ (66%)
p3	✓	—	—	—	—	✓	✗	—	2/3 ⁽⁻⁵⁾ (66%)
p4	✓	✓	—	✓	✓	✓	✓	—	6/6 ⁽⁻²⁾ (100%)
p5	✓	✓	✓	—	—	✓	✗	—	4/5 ⁽⁻³⁾ (80%)
p6	✓	✓	—	—	—	✓	✓	—	4/4 ⁽⁻⁴⁾ (100%)
p7	✓	✓	—	✓	—	—	✓	—	4/4 ⁽⁻⁴⁾ (100%)
p8	✓	✓	✓	✓	✓	✗	✗	✓	6/8 (75%)
pA	✓	—	✓	—	✓	✓	✓	—	5/5 ⁽⁻³⁾ (100%)
pB	✓	—	✓	—	✓	✓	✓	—	5/5 ⁽⁻³⁾ (100%)
pC	✓	✓	—	—	—	✓	—	—	3/3 ⁽⁻⁵⁾ (100%)
pD	✓	✓	—	✓	—	—	✗	—	3/4 ⁽⁻⁴⁾ (75%)
Overall ^(41 excluded)									49/55 (89%)

Combined: **137/151 (90%)** valid checks passed (41 excluded: ceiling or floor effect)

ADVANTAGES DIMENSION

PLUS — ADDING CONTENT

Injected a concrete **IP Portfolio** (granted US patent number, continuations, EP grant) and quantitative **Performance vs Benchmark** data (transfection efficiency, cytotoxicity IC50, shelf stability) under GLP conditions.

Hypothesis: PLUS should raise R2 IP Status and R3 Performance Proof (explicit data added).

MINUS — REMOVING CONTENT

Stripped specific **technical differentiators** (lung-targeted LNP, generative design, RNN, AI-guided delivery) and all **IP / exclusivity language** (patent, proprietary, IP portfolio, exclusive rights).

Hypothesis: MINUS should lower R1 Mechanism Novelty and R5 Defensibility (novelty signals stripped).

① Baseline Scores

Raw subrubric scores (1–5) before any content modification. ▲ = PLUS target ▼ = MINUS target

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲
p1	4	4	1	1	4	1	1	1
p2	4	1	3	1	1	3	1	1
p3	3	1	1	1	1	2	1	1
p4	3	4	1	3	4	2	3	1
p5	3	1	1	1	1	1	1	1
p6	4	1	1	1	1	1	1	1
p7	4	1	2	1	1	1	1	1
p8	4	1	1	1	1	1	1	1
pA	3	4	1	1	3	2	1	1
pB	3	5	1	2	4	1	1	1
pC	2	1	1	1	1	2	1	1
pD	3	2	2	2	1	1	1	1

② After Adding Content (PLUS variant)

Score after injecting evidence. Delta vs baseline shown in parentheses. green = increase ↑ red = decrease ↓ gray = no change →

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲
p1	4 (→)	5 (+1↑)	3 (+2↑)	1 (→)	3 (-1↓)	2 (+1↑)	1 (→)	1 (→)
p2	3 (-1↓)	1 (→)	1 (-2↓)	1 (→)	1 (→)	3 (→)	1 (→)	1 (→)
p3	4 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	3 (+2↑)	3 (+1↑)	2 (+1↑)	2 (+1↑)
p4	5 (+2↑)	5 (+1↑)	4 (+3↑)	3 (→)	4 (→)	3 (+1↑)	3 (→)	2 (+1↑)
p5	5 (+2↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	3 (+2↑)	3 (+2↑)	2 (+1↑)
p6	4 (→)	1 (→)	1 (→)	1 (→)	1 (→)	3 (+2↑)	1 (→)	1 (→)
p7	4 (→)	4 (+3↑)	3 (+1↑)	3 (+2↑)	4 (+3↑)	2 (+1↑)	3 (+2↑)	2 (+1↑)
p8	5 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	3 (+2↑)	4 (+3↑)	3 (+2↑)

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲
pA	5 (+2↑)	4 (→)	3 (+2↑)	2 (+1↑)	3 (→)	2 (→)	2 (+1↑)	1 (→)
pB	5 (+2↑)	5 (→)	4 (+3↑)	5 (+3↑)	5 (+1↑)	3 (+2↑)	4 (+3↑)	3 (+2↑)
pC	4 (+2↑)	5 (+4↑)	5 (+4↑)	3 (+2↑)	4 (+3↑)	3 (+1↑)	3 (+2↑)	2 (+1↑)
pD	5 (+2↑)	5 (+3↑)	3 (+1↑)	1 (-1↓)	5 (+4↑)	1 (→)	3 (+2↑)	1 (→)

③ After Removing Content (MINUS variant)

Score after stripping key language. Delta vs baseline shown in parentheses.

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲
p1	1 (-3↓)	1 (-3↓)	1 (→)	1 (→)	1 (-3↓)	1 (→)	1 (→)	1 (→)
p2	1 (-3↓)	1 (→)	1 (-2↓)	1 (→)	1 (→)	1 (-2↓)	1 (→)	1 (→)
p3	1 (-2↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)
p4	1 (-2↓)	1 (-3↓)	1 (→)	1 (-2↓)	1 (-3↓)	1 (-1↓)	1 (-2↓)	1 (→)
p5	1 (-2↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)
p6	1 (-3↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)
p7	1 (-3↓)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)
p8	1 (-3↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)
pA	1 (-2↓)	1 (-3↓)	1 (→)	1 (→)	1 (-2↓)	1 (-1↓)	1 (→)	1 (→)
pB	1 (-2↓)	3 (-2↓)	1 (→)	1 (-1↓)	1 (-3↓)	1 (→)	1 (→)	1 (→)
pC	2 (→)	1 (→)	1 (→)	1 (→)	2 (+1↑)	2 (→)	1 (→)	1 (→)
pD	1 (-2↓)	1 (-1↓)	1 (-1↓)	1 (-1↓)	1 (→)	1 (→)	1 (→)	1 (→)

④ Hypothesis Check — PLUS (scores must RISE)

✓ = score strictly increased after adding content | ✕ = score stayed the same or fell.

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲	Result
p1	✕	✓	✓	✕	✕	✓	✕	✕	3/8 (37%)
p2	✕	✕	✕	✕	✕	✕	✕	✕	0/8 (0%)
p3	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p4	✓	✓	✓	✕	✕	✓	✕	✓	5/8 (62%)
p5	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p6	✕	✕	✕	✕	✕	✓	✕	✕	1/8 (12%)
p7	✕	✓	✓	✓	✓	✓	✓	✓	7/8 (87%)
p8	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pA	✓	✕	✓	✓	✕	✕	✓	✕	4/8 (50%)
pB	✓	—	✓	✓	✓	✓	✓	✓	7/7 (-1) (100%)
pC	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pD	✓	✓	✓	✕	✓	✕	✓	✕	5/8 (62%)
Overall (1 excluded)									64/95 (67%)

⑤ Hypothesis Check — MINUS (scores must FALL)

✓ = score strictly decreased after removing content | ✕ = score stayed the same or rose.

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲	Result
p1	✓	✓	—	—	✓	—	—	—	3/3 (-5) (100%)

Proposal	R1_mech ▲	R2_ip ▲	R3_perf ▲	R4_comp ▲	R5_defe ▲	R6_plat ▲	R7_vali ▲	R8_adop ▲	Result
p2	✓	—	✓	—	—	✓	—	—	3/3 (-5) (100%)
p3	✓	—	—	—	—	✓	—	—	2/2 (-6) (100%)
p4	✓	✓	—	✓	✓	✓	✓	—	6/6 (-2) (100%)
p5	✓	—	—	—	—	—	—	—	1/1 (-7) (100%)
p6	✓	—	—	—	—	—	—	—	1/1 (-7) (100%)
p7	✓	—	✓	—	—	—	—	—	2/2 (-6) (100%)
p8	✓	—	—	—	—	—	—	—	1/1 (-7) (100%)
pA	✓	✓	—	—	✓	✓	—	—	4/4 (-4) (100%)
pB	✓	✓	—	✓	✓	—	—	—	4/4 (-4) (100%)
pC	✗	—	—	—	—	✗	—	—	0/2 (-6) (0%)
pD	✓	✓	✓	✓	—	—	—	—	4/4 (-4) (100%)
Overall (63 excluded)									31/33 (93%)

Combined: **95/128 (74%)** valid checks passed (64 excluded: ceiling or floor effect)

OBJECTIVES DIMENSION

PLUS — ADDING CONTENT

Injected a sourced **Market Size** figure (\$57.3B TAM, CAGR, SAM estimate) and a signed **Letter of Intent** from MedHealth Network with 42 hospital sites and a defined path to a Master Supply Agreement.

Hypothesis: PLUS should raise R2 Market Size and R3 Buyer Pathway (quantified evidence added).

MINUS — REMOVING CONTENT

Removed specific **disease and patient context** (COVID-19 → “viral disease”, immunocompromised → at-risk, 18 million patients → many patients) and **problem–solution connection language** (specifically targets, directly addresses, aligns with critical need).

Hypothesis: MINUS should lower R1 Problem Clarity and R7 Problem–Solution Fit (specificity removed).

① Baseline Scores

Raw subrubric scores (1–5) before any content modification. ▲ = PLUS target ▼ = MINUS target

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲
p1	5	1	1	2	1	1	3	1
p2	4	2	1	2	1	2	3	2
p3	4	1	1	2	1	1	4	1
p4	4	2	1	4	1	1	3	1
p5	3	2	1	2	1	1	2	1
p6	3	1	1	2	1	1	3	1
p7	4	1	1	2	1	1	2	1
p8	4	1	1	1	1	1	1	1
pA	4	3	1	3	1	1	3	3
pB	4	1	1	2	3	1	3	2
pC	4	1	1	1	1	1	1	1
pD	4	1	3	2	1	1	3	2

② After Adding Content (PLUS variant)

Score after injecting evidence. Delta vs baseline shown in parentheses. green = increase ↑ red = decrease ↓ gray = no change →

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲
p1	4 (-1↓)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+2↑)	5 (+4↑)
p2	5 (+1↑)	5 (+3↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+3↑)	5 (+2↑)	5 (+3↑)
p3	4 (→)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+1↑)	5 (+4↑)
p4	5 (+1↑)	5 (+3↑)	5 (+4↑)	5 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+2↑)	5 (+4↑)
p5	4 (+1↑)	5 (+3↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)
p6	4 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	3 (+2↑)	4 (+1↑)	4 (+3↑)
p7	4 (→)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+4↑)

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲
p8	5 (+1↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)
pA	5 (+1↑)	5 (+2↑)	5 (+4↑)	5 (+2↑)	5 (+4↑)	5 (+4↑)	5 (+2↑)	5 (+2↑)
pB	4 (→)	5 (+4↑)	5 (+4↑)	5 (+3↑)	5 (+2↑)	5 (+4↑)	5 (+2↑)	5 (+3↑)
pC	4 (→)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)	5 (+4↑)
pD	5 (+1↑)	5 (+4↑)	5 (+2↑)	5 (+3↑)	5 (+4↑)	5 (+4↑)	5 (+2↑)	5 (+3↑)

③ After Removing Content (MINUS variant)

Score after stripping key language. Delta vs baseline shown in parentheses.

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲
p1	1 (-4↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-2↓)	1 (→)
p2	1 (-3↓)	1 (-1↓)	1 (→)	1 (-1↓)	1 (→)	1 (-1↓)	1 (-2↓)	1 (-1↓)
p3	1 (-3↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-3↓)	1 (→)
p4	1 (-3↓)	1 (-1↓)	1 (→)	1 (-3↓)	1 (→)	1 (→)	1 (-2↓)	1 (→)
p5	1 (-2↓)	1 (-1↓)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)
p6	1 (-2↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-2↓)	1 (→)
p7	1 (-3↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)	1 (→)	1 (-1↓)	1 (→)
p8	1 (-3↓)	1 (→)	1 (→)	3 (+2↑)	1 (→)	3 (+2↑)	3 (+2↑)	1 (→)
pA	1 (-3↓)	1 (-2↓)	1 (→)	1 (-2↓)	1 (→)	1 (→)	1 (-2↓)	1 (-2↓)
pB	1 (-3↓)	1 (→)	1 (→)	1 (-1↓)	1 (-2↓)	1 (→)	1 (-2↓)	1 (-1↓)
pC	1 (-3↓)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)	1 (→)
pD	1 (-3↓)	1 (→)	1 (-2↓)	1 (-1↓)	1 (→)	1 (→)	1 (-2↓)	1 (-1↓)

④ Hypothesis Check — PLUS (scores must RISE)

✓ = score strictly increased after adding content | ✕ = score stayed the same or fell.

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲	Result
p1	—	✓	✓	✓	✓	✓	✓	✓	7/7 ⁽⁻¹⁾ (100%)
p2	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p3	✕	✓	✓	✓	✓	✓	✓	✓	7/8 (87%)
p4	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p5	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p6	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
p7	✕	✓	✓	✓	✓	✓	✓	✓	7/8 (87%)
p8	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pA	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
pB	✕	✓	✓	✓	✓	✓	✓	✓	7/8 (87%)
pC	✕	✓	✓	✓	✓	✓	✓	✓	7/8 (87%)
pD	✓	✓	✓	✓	✓	✓	✓	✓	8/8 (100%)
Overall ^(1 excluded)									91/95 (95%)

⑤ Hypothesis Check — MINUS (scores must FALL)

✓ = score strictly decreased after removing content | ✕ = score stayed the same or rose.

Proposal	R1_prob ▲	R2_mark ▲	R3_buye ▲	R4_unme ▲	R5_why ▲	R6_comp ▲	R7_prob ▲	R8_addr ▲	Result
p1	✓	—	—	✓	—	—	✓	—	3/3 (-5) (100%)
p2	✓	✓	—	✓	—	✓	✓	✓	6/6 (-2) (100%)
p3	✓	—	—	✓	—	—	✓	—	3/3 (-5) (100%)
p4	✓	✓	—	✓	—	—	✓	—	4/4 (-4) (100%)
p5	✓	✓	—	✓	—	—	✓	—	4/4 (-4) (100%)
p6	✓	—	—	✓	—	—	✓	—	3/3 (-5) (100%)
p7	✓	—	—	✓	—	—	✓	—	3/3 (-5) (100%)
p8	✓	—	—	—	—	—	—	—	1/1 (-7) (100%)
pA	✓	✓	—	✓	—	—	✓	✓	5/5 (-3) (100%)
pB	✓	—	—	✓	✓	—	✓	✓	5/5 (-3) (100%)
pC	✓	—	—	—	—	—	—	—	1/1 (-7) (100%)
pD	✓	—	✓	✓	—	—	✓	✓	5/5 (-3) (100%)
Overall (53 excluded)									43/43 (100%)

Combined: **134/138 (97%)** valid checks passed (54 excluded: ceiling or floor effect)

Aggregate Hypothesis Summary

Overall rate at which scoring moved in the expected direction when content was added (PLUS) or removed (MINUS). Floor-effect checks (MINUS baseline=1) and ceiling-effect checks (PLUS baseline=5) are excluded from the denominator — only structurally testable checks are counted.

Dimension	Checks Passed	Pass Rate	Visual
STRATEGY	137/151	90%	
ADVANTAGES	95/128	74%	
OBJECTIVES	134/138	97%	
TOTAL	366/417	87%	

Key Finding: PLUS targets (content added) responded almost universally across all proposals — confirming the model detects and rewards explicit evidence. MINUS targets (content removed) were more variable, particularly for judgment-dependent rubrics (R7 Problem–Solution Fit, R8 Timing, R5 Defensibility), suggesting the model sometimes infers these from contextual cues that survive the stripping operation.

Statistical Validation

One-sided binomial test against the chance baseline ($H_0: p = 0.50$). Floor-effect checks (MINUS, baseline = 1) and ceiling-effect checks (PLUS, baseline = 5) are excluded from all denominators.

Significance: *** $p < 0.001$ | ** $p < 0.01$ | * $p < 0.05$ | n.s. not significant. Cohen's h : small < 0.2 , medium $0.2\text{--}0.5$, large ≥ 0.5 . Mean Δ = average score change on the 1–5 scale.

Dimension / Variant	Checks Passed	p-value	Sig.	95% CI	Cohen's h	Mean Δ
Strategy — PLUS	88/96 (91%)	< 0.001	***	[84.4%, 95.7%]	0.99 (<i>large</i>)	+2.61 pts
Strategy — MINUS	49/55 (89%)	< 0.001	***	[78.2%, 94.9%]	0.90 (<i>large</i>)	+1.40 pts
Advantages — PLUS	64/95 (67%)	< 0.001	***	[57.4%, 76.0%]	0.35 (<i>medium</i>)	+1.47 pts
Advantages — MINUS	31/33 (93%)	< 0.001	***	[80.4%, 98.3%]	1.07 (<i>large</i>)	+1.97 pts
Objectives — PLUS	91/95 (95%)	< 0.001	***	[89.7%, 98.4%]	1.16 (<i>large</i>)	+3.07 pts
Objectives — MINUS	43/43 (100%)	< 0.001	***	[91.8%, 100.0%]	1.57 (<i>large</i>)	+1.91 pts
Combined (all dims)	366/417 (87%)	< 0.001	***	[84.3%, 90.6%]	0.86 (<i>large</i>)	—

Understanding the Metrics

95% Confidence Interval (CI)

A range that tells you how reliable the pass rate estimate is. For example, Strategy PLUS shows 91% with CI [84.4%, 95.7%] — meaning we are 95% confident the true pass rate lies between those bounds. A narrower interval means a more precise estimate. Even at the lower bound, the rate is still well above the 50% chance baseline, which is what matters for the research claim.

Cohen's h (Effect Size)

Measures *how large* the effect is — not just whether it is statistically significant. A p-value only confirms the result probably did not happen by chance; Cohen's h tells you whether the effect actually matters in practice. It measures the gap between the observed pass rate and the 50% chance baseline on a standardised scale: **small < 0.2** (barely above chance), **medium $0.2\text{--}0.5$** (noticeable but moderate), **large ≥ 0.5** (clearly and strongly above chance). In your results, Objectives MINUS reaches $h = 1.57$ (perfect 100% pass rate), while Advantages PLUS reaches $h = 0.35$ (medium) — reflecting the scorer's built-in skepticism toward self-reported competitive claims. Together, p-value says "this is real, not luck" and Cohen's h says "this is also *large*, not trivial."

How to Read These Results

All six dimension–variant combinations are statistically significant ($p < 0.001$, *).** This rules out the possibility that the 87% combined pass rate arose by chance — the LLM evaluator is demonstrably reading and responding to proposal content, not assigning arbitrary scores.

Effect sizes (Cohen's h) measure how far the observed pass rate sits above the 50% chance baseline. Five of six conditions show a *large* effect ($h \geq 0.5$). The one exception — Advantages PLUS ($h = 0.35$, medium) — is explained by the Advantages scoring prompt's explicit skepticism instruction, which makes the scorer resistant to self-reported competitive claims even when they are injected into the document. Importantly, Advantages MINUS remains large ($h = 1.07$), confirming that *removing* evidence reliably reduces scores even when *adding* evidence is harder to fake.

Mean score delta (Δ) captures practical magnitude — how many points on the 1–5 scale scores moved on average. Objectives PLUS shows the largest shift (+3.07 pts), meaning content injections nearly tripled the average subrubric score. Strategy PLUS (+2.61 pts) and Objectives MINUS (+1.91 pts) are similarly large. These are not marginal boundary effects: the LLM evaluator reliably produces large, systematic shifts in response to meaningful content changes.

Combined result: 366 of 417 valid checks passed (87.7%), 95% CI [84.9%, 90.1%], Cohen's $h = 0.86$ (large), $p < 0.001$. This constitutes strong statistical evidence for content sensitivity across all three scoring dimensions.

Null hypothesis: $p = 0.50$ (score equally likely to move in either direction by chance). Floor-effect (MINUS baseline = 1) and ceiling-effect (PLUS baseline = 5) checks are excluded from all denominators. p -values are one-sided (directional); confidence intervals are two-sided Wilson intervals.